

# TMDB Dataset Investigation

November 16, 2018

```
# Investigate TMDB Dataset
```

## 0.1 Table of Contents

Introduction

    Data Wrangling

    Exploratory Data Analysis

```
<ul>
    <li><a href="#q1">What are the key trends of the movie industry over the year?</a></li>
    <li><a href="#q2">What kind of properties associated with high-profit movies?</a></li>
</ul>
```

Conclusions

## Introduction

This report will try to highlight trends & correlations observed in TMDB dataset by answering the following questions,

- What are the key trends of the movie industry over the year?
- What kind of properties associated with high-profit movies?

```
In [1]: # import necessary packages for analysis
```

```
    import numpy as np
    import pandas as pd
    import matplotlib.pyplot as plt
    import seaborn as sns
```

```
    # Configure initial settings for visualization
    sns.set()
```

```
% matplotlib inline
```

```
## Data Wrangling
```

### 0.1.1 General Properties

```
In [2]: # Load your data and print out a few lines. Perform operations to inspect data
#   types and look for instances of missing or possibly errant data.

# Load / inspect data
df = pd.read_csv('tmdb-movies.csv')
df.head(2)

Out[2]:      id    imdb_id  popularity      budget      revenue  original_title \
0  135397  tt0369610    32.985763  150000000  1513528810        Jurassic World
1  76341   tt1392190    28.419936  150000000  378436354  Mad Max: Fury Road

                                         cast \
0  Chris Pratt|Bryce Dallas Howard|Irrfan Khan|Vi...
1  Tom Hardy|Charlize Theron|Hugh Keays-Byrne|Nic...

                                         homepage      director      tagline \
0  http://www.jurassicworld.com/  Colin Trevorrow  The park is open.
1  http://www.madmaxmovie.com/  George Miller  What a Lovely Day.

                                         overview runtime \
0  ...          Twenty-two years after the events of Jurassic ...       124
1  ...          An apocalyptic story set in the furthest reach...       120

                                         genres \
0  Action|Adventure|Science Fiction|Thriller
1  Action|Adventure|Science Fiction|Thriller

                                         production_companies release_date vote_count \
0  Universal Studios|Amblin Entertainment|Legenda...       6/9/15      5562
1  Village Roadshow Pictures|Kennedy Miller Produ...      5/13/15      6185

      vote_average  release_year      budget_adj      revenue_adj
0            6.5        2015  1.379999e+08  1.392446e+09
1            7.1        2015  1.379999e+08  3.481613e+08

[2 rows x 21 columns]
```

```
In [3]: df.shape
```

```
Out[3]: (10866, 21)
```

```
In [4]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10866 entries, 0 to 10865
Data columns (total 21 columns):
id                10866 non-null int64
```

```

imdb_id           10856 non-null object
popularity        10866 non-null float64
budget            10866 non-null int64
revenue           10866 non-null int64
original_title    10866 non-null object
cast              10790 non-null object
homepage          2936 non-null object
director          10822 non-null object
tagline           8042 non-null object
keywords          9373 non-null object
overview          10862 non-null object
runtime           10866 non-null int64
genres            10843 non-null object
production_companies 9836 non-null object
release_date      10866 non-null object
vote_count         10866 non-null int64
vote_average       10866 non-null float64
release_year       10866 non-null int64
budget_adj         10866 non-null float64
revenue_adj        10866 non-null float64
dtypes: float64(4), int64(6), object(11)
memory usage: 1.7+ MB

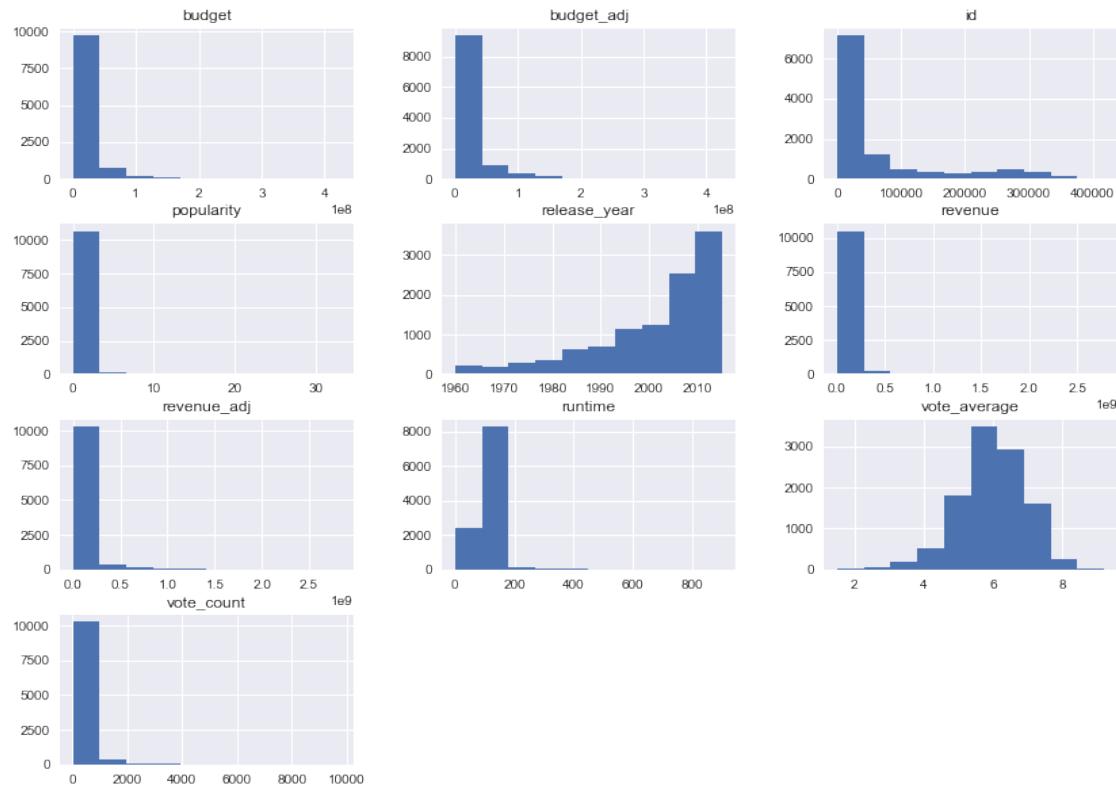
```

Identifying properties associated with different columns (min, max, percental distribution) before cleaning.

In [5]: df.describe()

	id	popularity	budget	revenue	runtime	\
count	10866.000000	10866.000000	1.086600e+04	1.086600e+04	10866.000000	
mean	66064.177434	0.646441	1.462570e+07	3.982332e+07	102.070863	
std	92130.136561	1.000185	3.091321e+07	1.170035e+08	31.381405	
min	5.000000	0.000065	0.000000e+00	0.000000e+00	0.000000	
25%	10596.250000	0.207583	0.000000e+00	0.000000e+00	90.000000	
50%	20669.000000	0.383856	0.000000e+00	0.000000e+00	99.000000	
75%	75610.000000	0.713817	1.500000e+07	2.400000e+07	111.000000	
max	417859.000000	32.985763	4.250000e+08	2.781506e+09	900.000000	
	vote_count	vote_average	release_year	budget_adj	revenue_adj	
count	10866.000000	10866.000000	10866.000000	1.086600e+04	1.086600e+04	
mean	217.389748	5.974922	2001.322658	1.755104e+07	5.136436e+07	
std	575.619058	0.935142	12.812941	3.430616e+07	1.446325e+08	
min	10.000000	1.500000	1960.000000	0.000000e+00	0.000000e+00	
25%	17.000000	5.400000	1995.000000	0.000000e+00	0.000000e+00	
50%	38.000000	6.000000	2006.000000	0.000000e+00	0.000000e+00	
75%	145.750000	6.600000	2011.000000	2.085325e+07	3.369710e+07	
max	9767.000000	9.200000	2015.000000	4.250000e+08	2.827124e+09	

```
In [6]: #visualize different columns in the data frame using histograms.  
df.hist(figsize=(14,10));
```



There are a significant count of movies with `revenue`, `revenue_adj`, `budget`, and `budget_adj` == zero.

```
In [7]: # Countifying amount of entries with zero value in `budget_adj` or `revenue_adj`  
(df['budget_adj']==0).sum(), (df['revenue_adj']==0).sum()
```

```
Out[7]: (5696, 6016)
```

There are a significant count of movies with `runtime` == zero.

```
In [8]: # Countifying amount of entries with zero value in `runtime`  
(df['runtime']==0).sum()
```

```
Out[8]: 31
```

Based on initial data exploration, the following actions to be taken on different columns

- **Columns to be dropped for irrelevance to investigation objectives**
  - `id`, `cast`, `homepage`, `keywords`, `director`, `tagline`, `overview`, `production_companies`
- **Columns to be dropped for redundancy**

- budget: Already covered under the budget\_adj column
- revenue: Already covered under the revenue\_adj column
- **Columns to be created, manipulated, changed or repurposed**
  - imdb\_id to be used as an index (will be useful as a unique identifier of different records).
  - release\_date will be dropped after been used to create a new column release\_month.
  - budget\_adj and revenue\_adj have values equals 0. We will change those to NaN to make sure they won't affect mathematical calculation in later stages. > We won't fill budget\_adj and revenue\_adj zero items with average because we have a significant amount of entries with (~5k entries) and filling them with average will skew analysis associated with revenue-budget correlation.
  - runtime have items == 0 (30 items). We will fill those with average.
  - Create a new column profit\_adj = revenue\_adj - budget\_adj

## 0.1.2 Data Cleaning

### Dropping redundant / irrelevant columns

```
In [9]: # dropping irrelevant columns
df.drop(['id', 'cast', 'homepage', 'keywords', 'director', 'tagline', 'overview', 'production_c
          # dropping redundant columns
          df.drop(['budget', 'revenue'], axis=1, inplace=True)
```

```
In [10]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10866 entries, 0 to 10865
Data columns (total 11 columns):
imdb_id           10856 non-null object
popularity        10866 non-null float64
original_title    10866 non-null object
runtime           10866 non-null int64
genres            10843 non-null object
release_date      10866 non-null object
vote_count         10866 non-null int64
vote_average      10866 non-null float64
release_year      10866 non-null int64
budget_adj        10866 non-null float64
revenue_adj        10866 non-null float64
dtypes: float64(4), int64(3), object(4)
memory usage: 933.9+ KB
```

### Detect & remove duplicate rows

```
In [11]: df.duplicated().sum()
```

```
Out[11]: 1
```

```
In [12]: df.drop_duplicates(inplace=True)
```

## Detect & fill/removeNaN values

In [13]: `df.isnull().sum()`

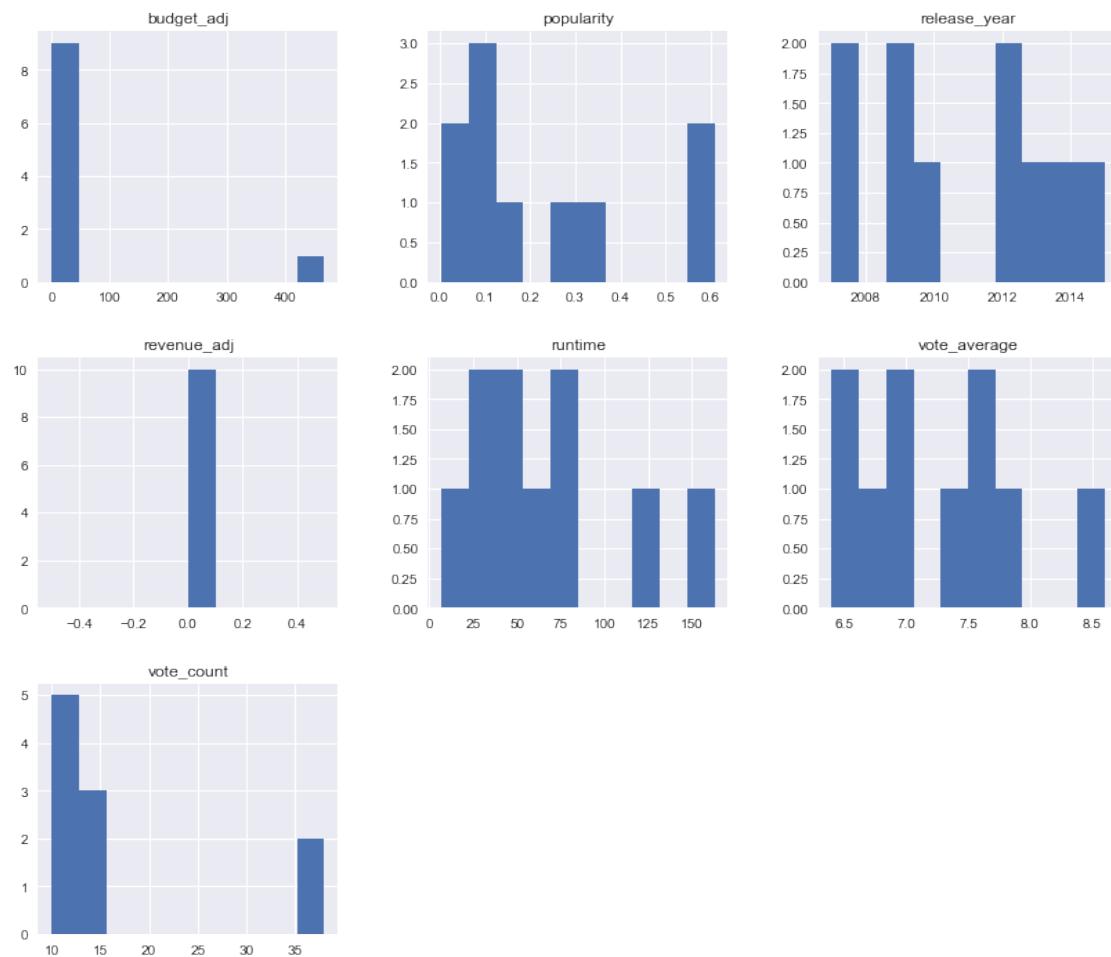
Out[13]:

imdb_id	10
popularity	0
original_title	0
runtime	0
genres	23
release_date	0
vote_count	0
vote_average	0
release_year	0
budget_adj	0
revenue_adj	0

dtype: int64

### Nan in imdb\_id column

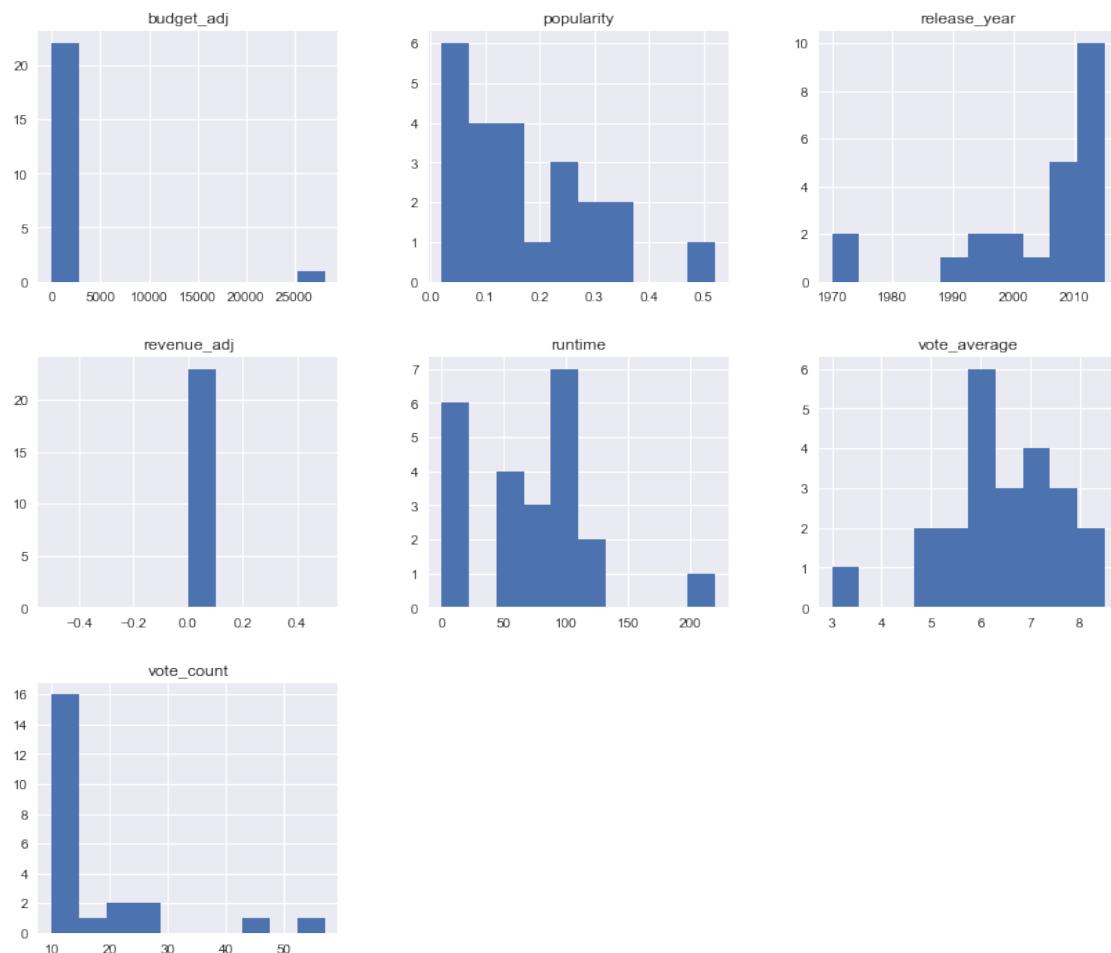
In [14]: `df[df.imdb_id.isnull()].hist(figsize=(14,12));`



We found that all items with NaN `imdb_id` have zero value in both `budget_adj` and `revenue_adj`. Besides that, we found no other correlation with other properties in the datasets. Hence, it will be ok if we drop the 10 records of NaN in `imdb_id`.

### **NaN in genres column**

```
In [15]: df[df.genres.isnull()].hist(figsize=(14,12));
```



We found that all items with NaN `genres` have zero value in both `budget_adj` and `revenue_adj`. Besides that, we found no other correlation with other properties in the datasets. Hence, it will be ok if we drop the 23 records of NaN in `genres`.

### **Dropping NaN rows**

```
In [16]: df.dropna(inplace = True)
df.isnull().sum()
```

```
Out[16]: imdb_id      0
popularity      0
```

```
original_title      0
runtime            0
genres             0
release_date       0
vote_count         0
vote_average       0
release_year       0
budget_adj         0
revenue_adj        0
dtype: int64
```

**Manipulating columns for analysis** Create the `release_month` from `release_year` then drop `release_year`.

```
In [17]: df['release_month'] = pd.DatetimeIndex(df['release_date']).month
df.drop(['release_date'], axis=1, inplace=True)
```

Fill into zero items in `budget_adj` and `revenue_adj` with NaN

```
In [18]: # Fill into zero items in budget_adj and revenue_adj with NaN
df['budget_adj'] = df['budget_adj'].replace(0,np.NAN)
df['revenue_adj'] = df['revenue_adj'].replace(0,np.NAN)
```

Fill into zero items in `runtime` with average

```
In [19]: # Fill into zero items in budget_adj and revenue_adj with NaN
df['runtime'] = df['runtime'].replace(0,np.NAN)
df['runtime'].fillna(df.runtime.mean(), inplace = True)
```

**Create new column profit\_adj**

```
In [20]: #create new column `profit_adj`
df['profit_adj'] = df['revenue_adj'] - df['budget_adj']
```

Previous step concludes the data wrangling phase. Below is the general structure and properties of the cleaned data-set.

```
In [21]: df.isnull().sum()
```

```
Out[21]: imdb_id          0
popularity        0
original_title    0
runtime           0
genres            0
vote_count        0
vote_average      0
release_year      0
budget_adj        5667
revenue_adj       5985
release_month     0
profit_adj        6980
dtype: int64
```

```
In [22]: df.info()
```

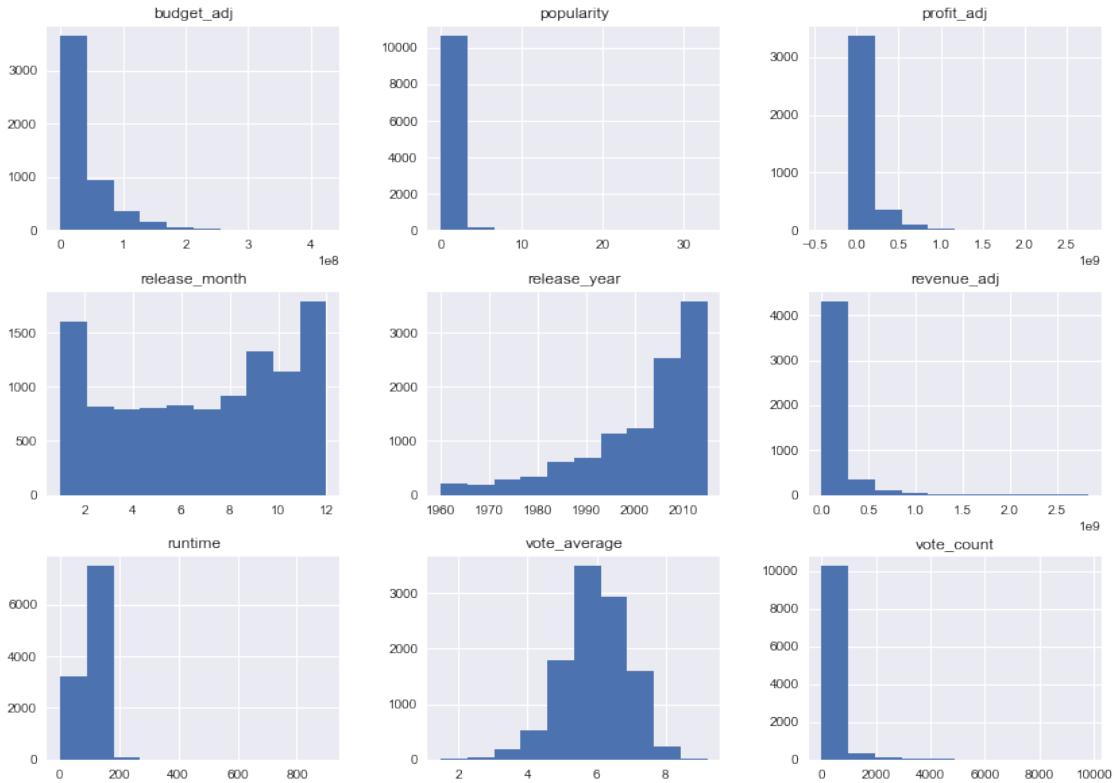
```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 10834 entries, 0 to 10865
Data columns (total 12 columns):
imdb_id           10834 non-null object
popularity        10834 non-null float64
original_title    10834 non-null object
runtime           10834 non-null float64
genres            10834 non-null object
vote_count         10834 non-null int64
vote_average      10834 non-null float64
release_year       10834 non-null int64
budget_adj         5167 non-null float64
revenue_adj        4849 non-null float64
release_month     10834 non-null int64
profit_adj         3854 non-null float64
dtypes: float64(6), int64(3), object(3)
memory usage: 1.1+ MB
```

```
In [23]: df.describe()
```

```
Out[23]:   popularity      runtime    vote_count  vote_average  release_year \
count  10834.000000  10834.000000  10834.000000  10834.000000  10834.000000
mean    0.647776    102.446409   217.970833    5.973159    2001.308196
std     1.001316    30.798047   576.368151    0.933831    12.815839
min     0.000065    2.000000    10.000000    1.500000    1960.000000
25%    0.208387    90.000000   17.000000    5.400000    1995.000000
50%    0.384587    99.000000   38.000000    6.000000    2006.000000
75%    0.715767   111.000000  146.000000    6.600000    2011.000000
max    32.985763   900.000000  9767.000000   9.200000    2015.000000

                                budget_adj  revenue_adj  release_month  profit_adj
count  5.167000e+03  4.849000e+03  10834.000000  3.854000e+03
mean   3.690334e+07  1.151009e+08   6.829149  9.282470e+07
std    4.196281e+07  1.988557e+08   3.439508  1.940715e+08
min    9.210911e-01  2.370705e+00   1.000000 -4.139124e+08
25%   8.102293e+06  1.046585e+07   4.000000 -1.504995e+06
50%   2.273036e+07  4.395666e+07   7.000000  2.737064e+07
75%   5.008384e+07  1.316482e+08  10.000000  1.074548e+08
max   4.250000e+08  2.827124e+09  12.000000  2.750137e+09
```

```
In [24]: #visualize different columns in the data frame using histograms.
df.hist(figsize=(14,10));
```



## ## Exploratory Data Analysis

This investigation will try to address the following questions

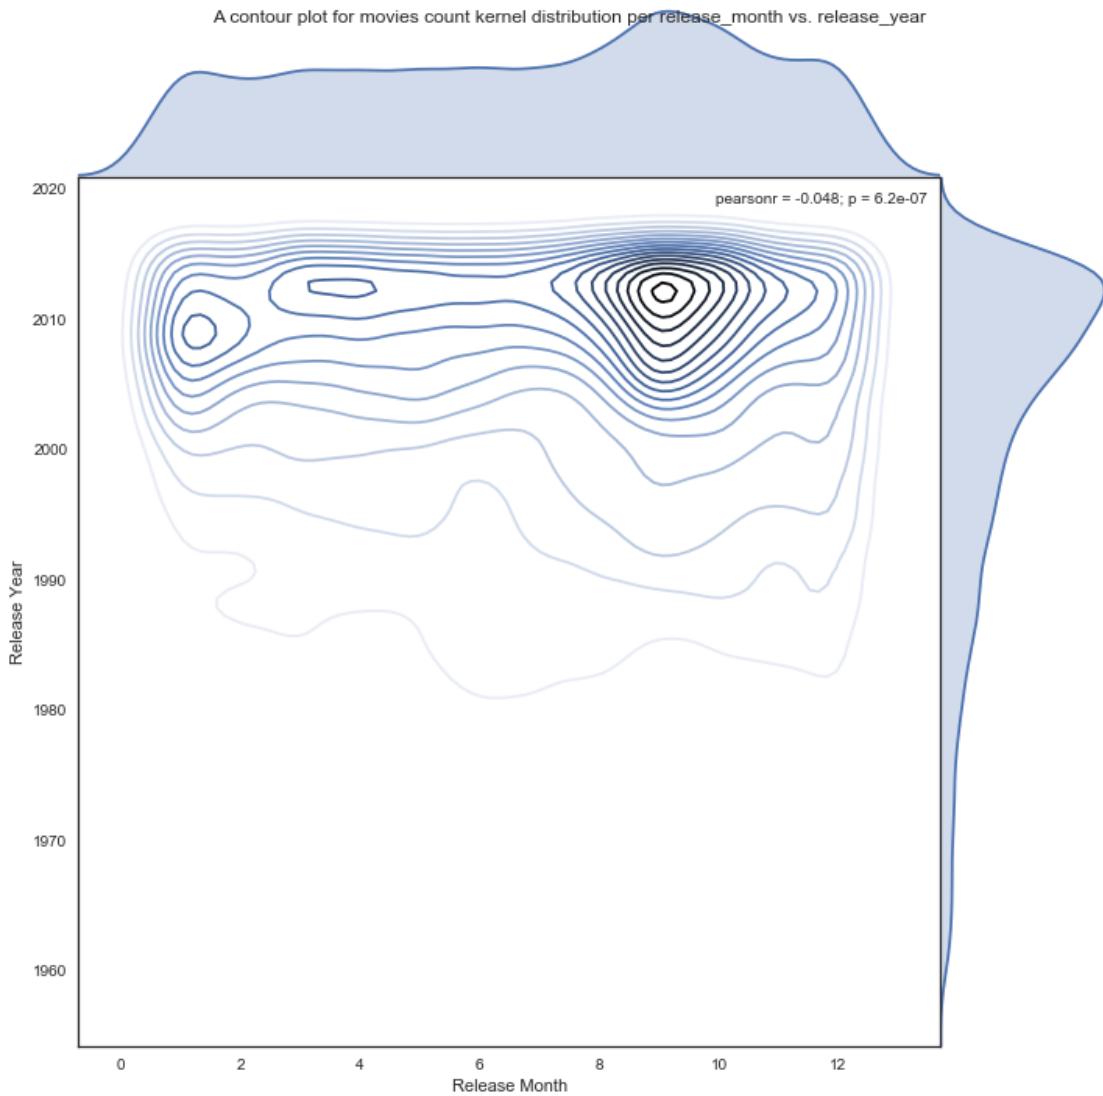
- What are the key trends of the movie industry over the year?
- What kind of properties associated with high-profit movies?

### What are the key trends of the movie industry over the year? This investigation will focus on trends associated with **Volume**, **Popularity**, **Run Time** and **Release time**.

Analyzing the movies release volume over release year & release month.

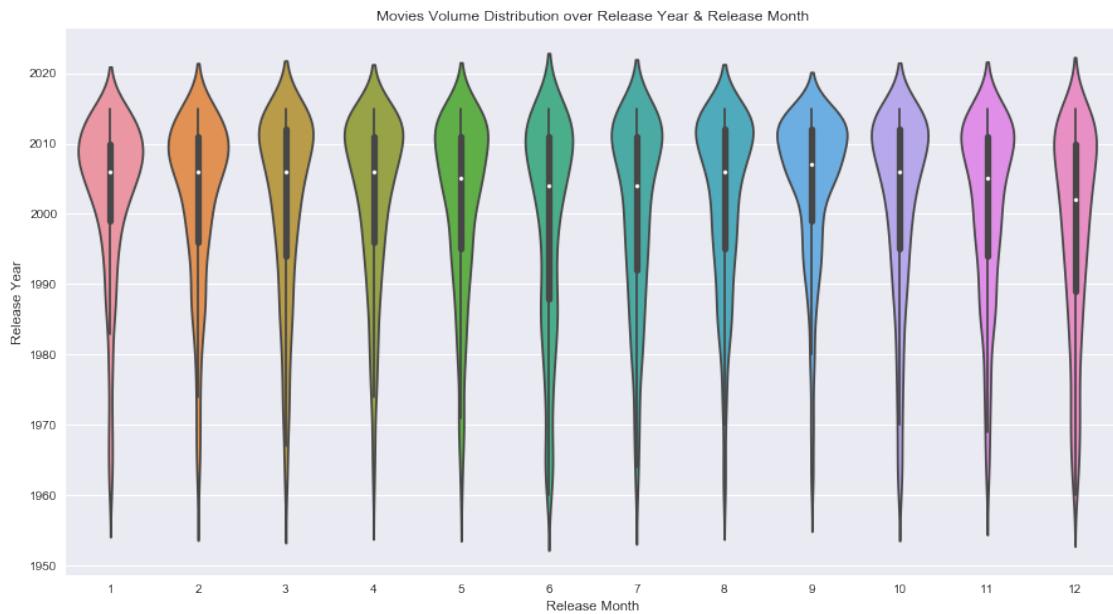
```
In [25]: # Create a contour plot for release_month vs. release_year
with sns.axes_style('white'):
    ax = sns.jointplot("release_month", "release_year", df, kind='kde',
                       size=10, space=0, shade=False, n_levels=20);
    ax.ax_joint.set_xlabel('Release Month')
    ax.ax_joint.set_ylabel('Release Year')
    ax.ax_joint.set_label('Helloworld')

    ax.fig.suptitle('A contour plot for movies count kernel distribution per release_mo
```

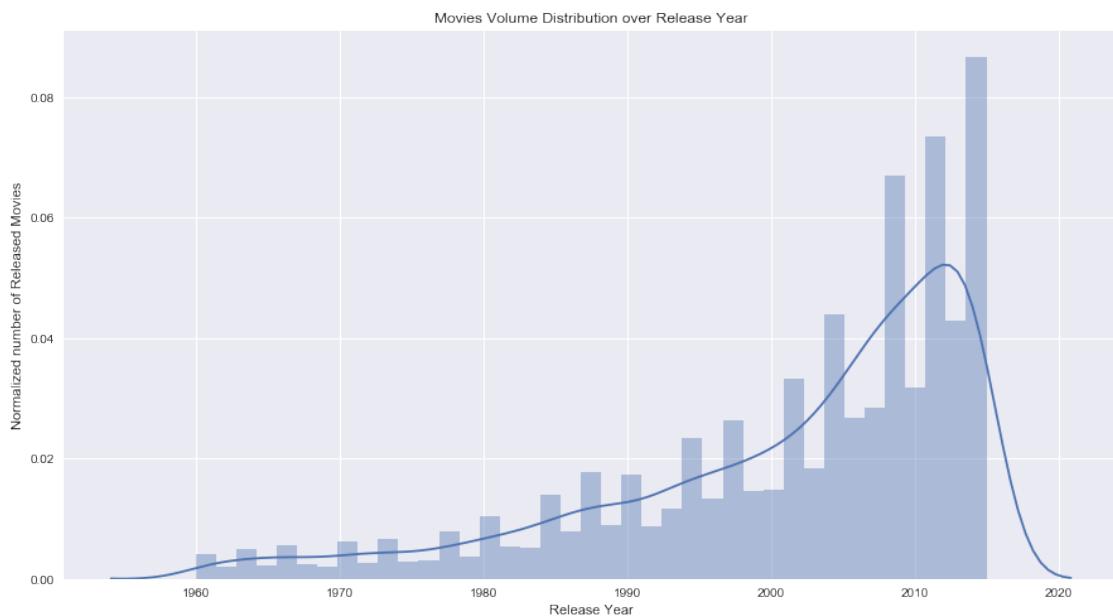


Note: We didn't use `xlim` and `ylim` on above curve to have better visualization, but it has to be clear that the dataset at hand holds data for the period from **1960** to **2015** and tail drop appearing on the graph is just an **interpolation** by the plotting tool.

```
In [26]: # Create a violin plot for release_month vs. release_year
plt.figure(figsize=(15,8))
ax = sns.violinplot("release_month", "release_year", data=df, );
plt.title('Movies Volume Distribution over Release Year & Release Month')
plt.xlabel('Release Month')
plt.ylabel('Release Year');
```



```
In [27]: # Create a normalized distribution histogram with KDE associated with volume trends over time
plt.figure(figsize=(15,8))
sns.distplot(df['release_year'], kde=True, norm_hist=True)
plt.title('Movies Volume Distribution over Release Year')
plt.xlabel('Release Year')
plt.ylabel('Normalized number of Released Movies');
```



A Couple of observations from the above diagrams, - This dataset infers a significant increase in the volume of movies over the years. - A pattern change has happened to release month around the 80s-years that require further analysis.

Investigating the volume increase over the years.

```
In [28]: df['release_year'].describe()
```

```
Out[28]: count    10834.000000
          mean     2001.308196
          std      12.815839
          min      1960.000000
          25%     1995.000000
          50%     2006.000000
          75%     2011.000000
          max      2015.000000
          Name: release_year, dtype: float64
```

Above IQR of `release_year` states that half of movies in this data-set have been released after 2006 (50% point).

Above diagram highlights two observations - Significant increase in released movies per year in the last 15 years (2000 till 2015). - Movies tend to be released in bursts every couple of years.

### Investigating the pattern change of release month over years

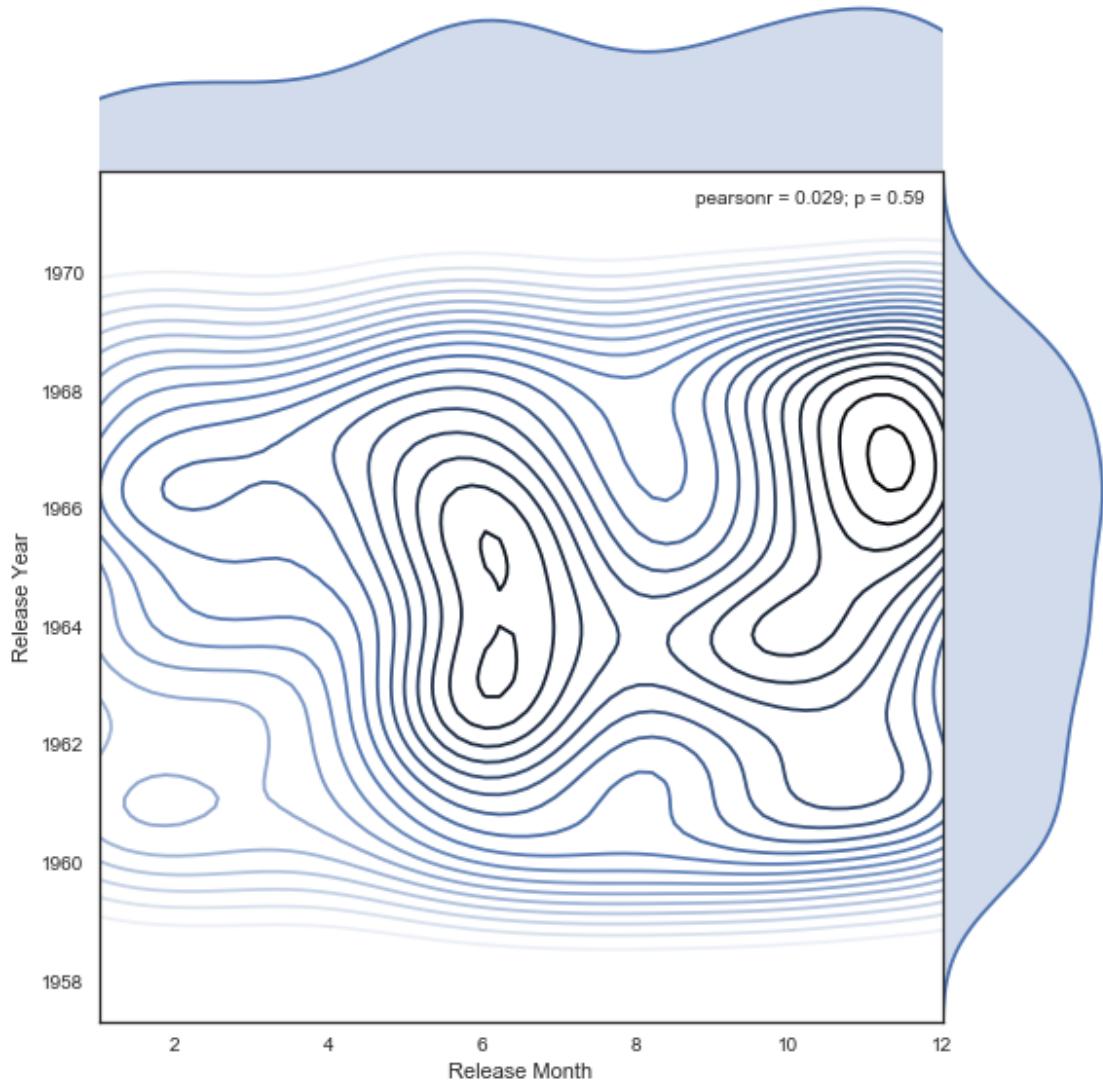
Create a categorical column `release_decade` from `release_year` around its x10 increments (ex. 1960s, 1970s, etc.)

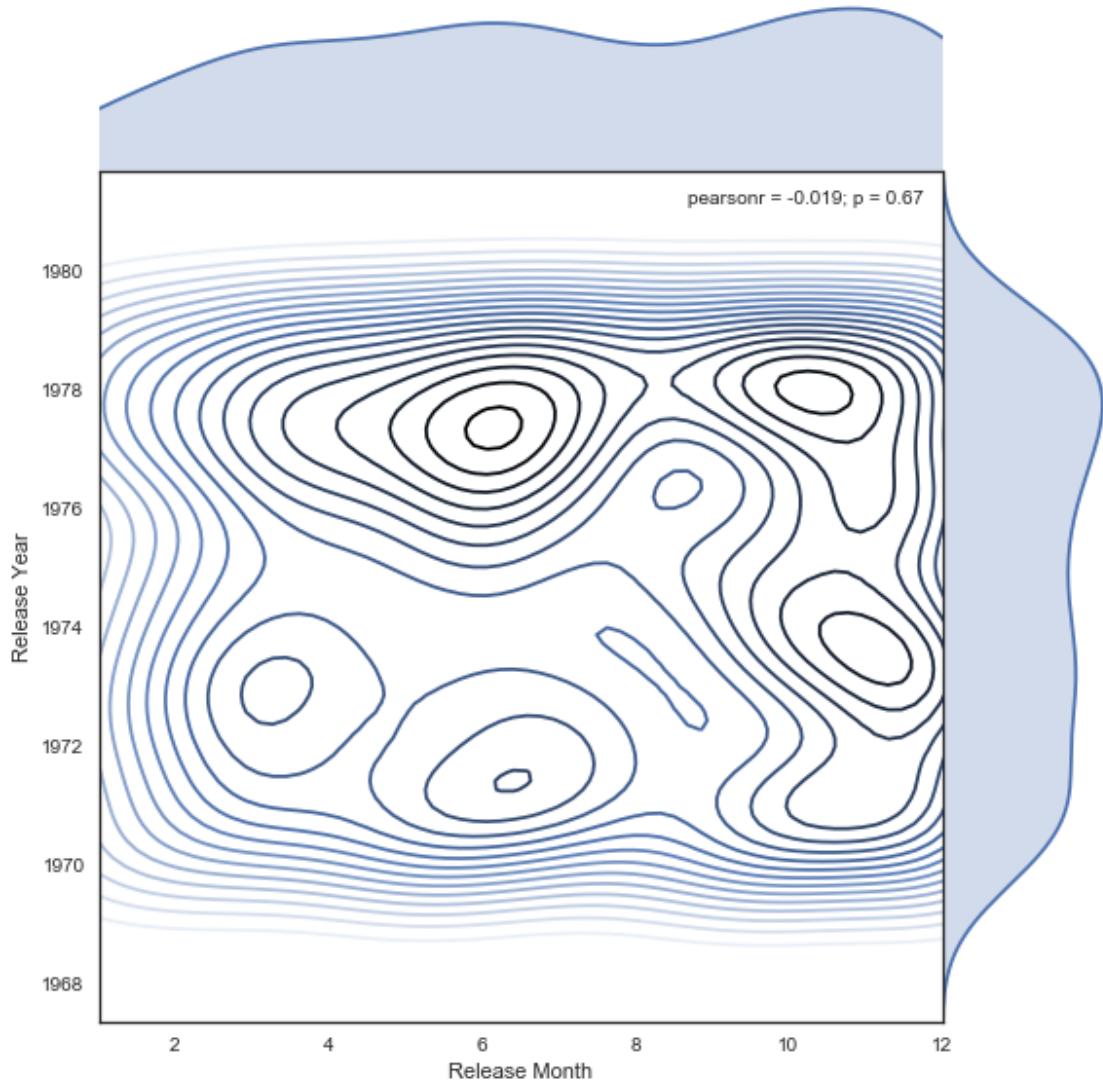
```
In [29]: # Create release_decade categorical column
bin_edges = [1959, 1969, 1979, 1989, 1999, 2009, 2019]
bin_names = ['1960s', '1970s', '1980s', '1990s', '2000s', '2010s']
df['release_decade'] = pd.cut(df['release_year'], bin_edges, labels=bin_names)
df['release_decade'].value_counts()
```

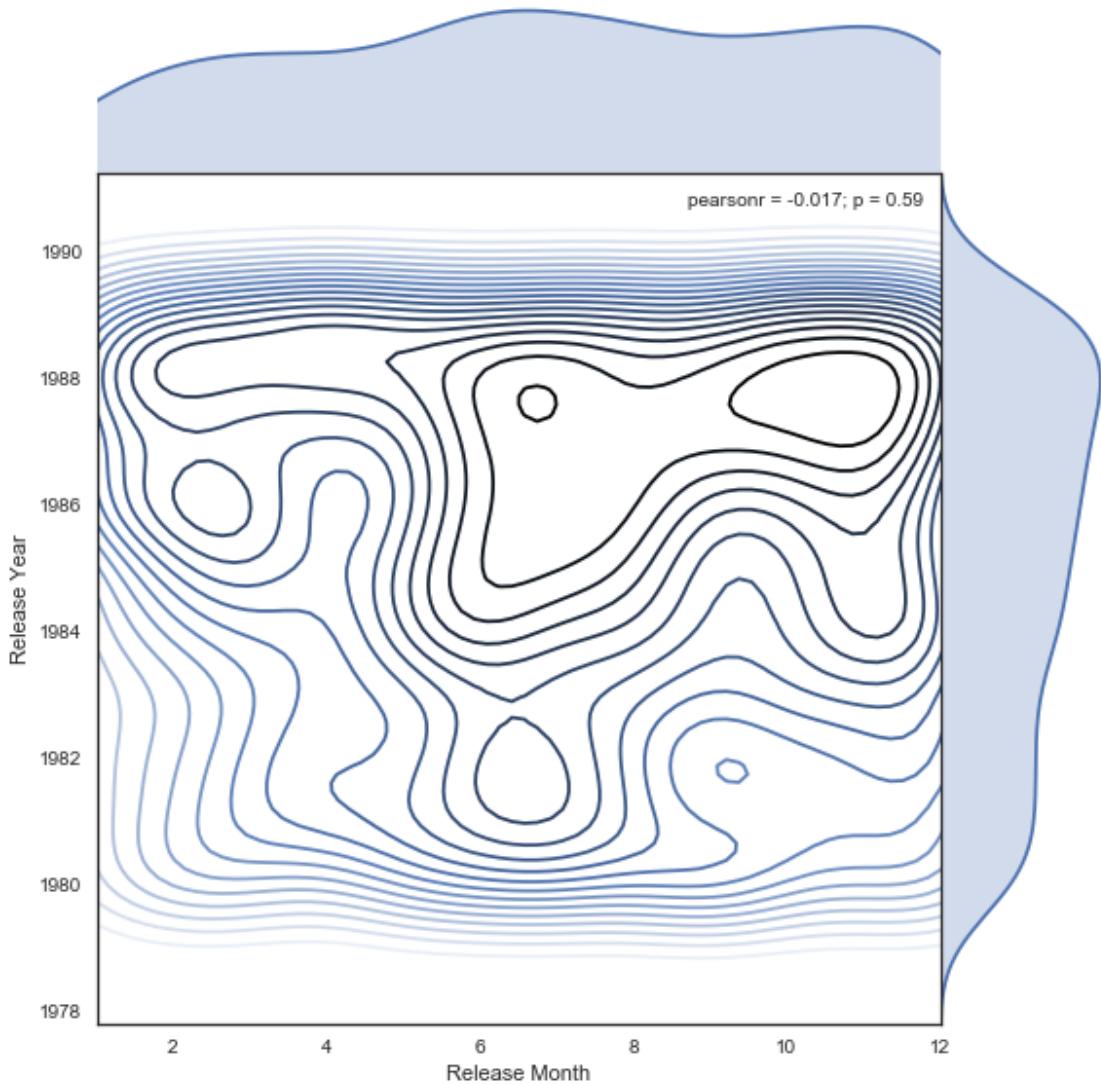
```
Out[29]: 2010s    3589
          2000s    3552
          1990s    1763
          1980s    1062
          1970s     506
          1960s     362
          Name: release_decade, dtype: int64
```

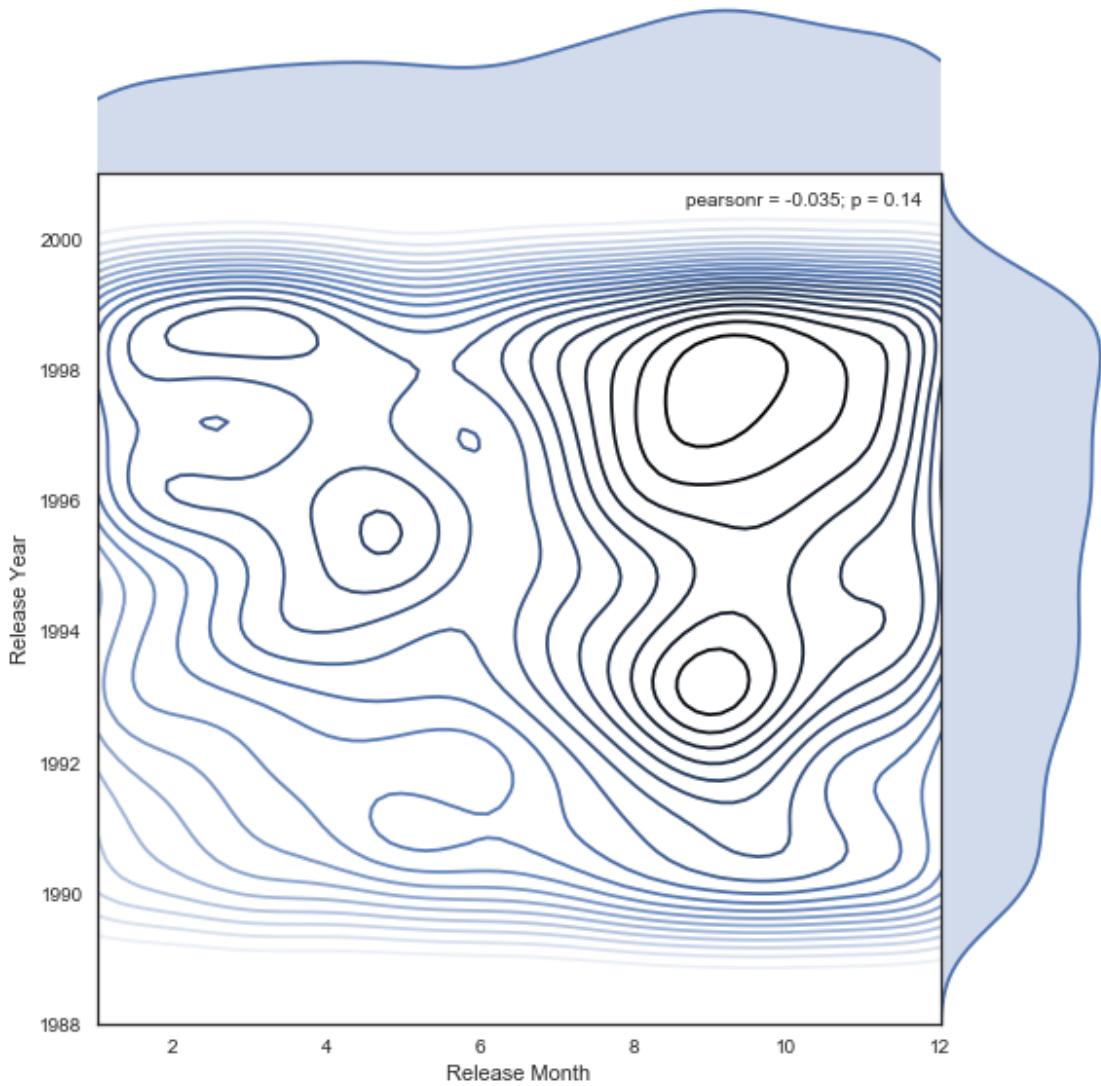
Create a counter plot with kernel distribution estimation (kde) graph of movie-count grouped by `decade` vs. `release_year` and `release_month`

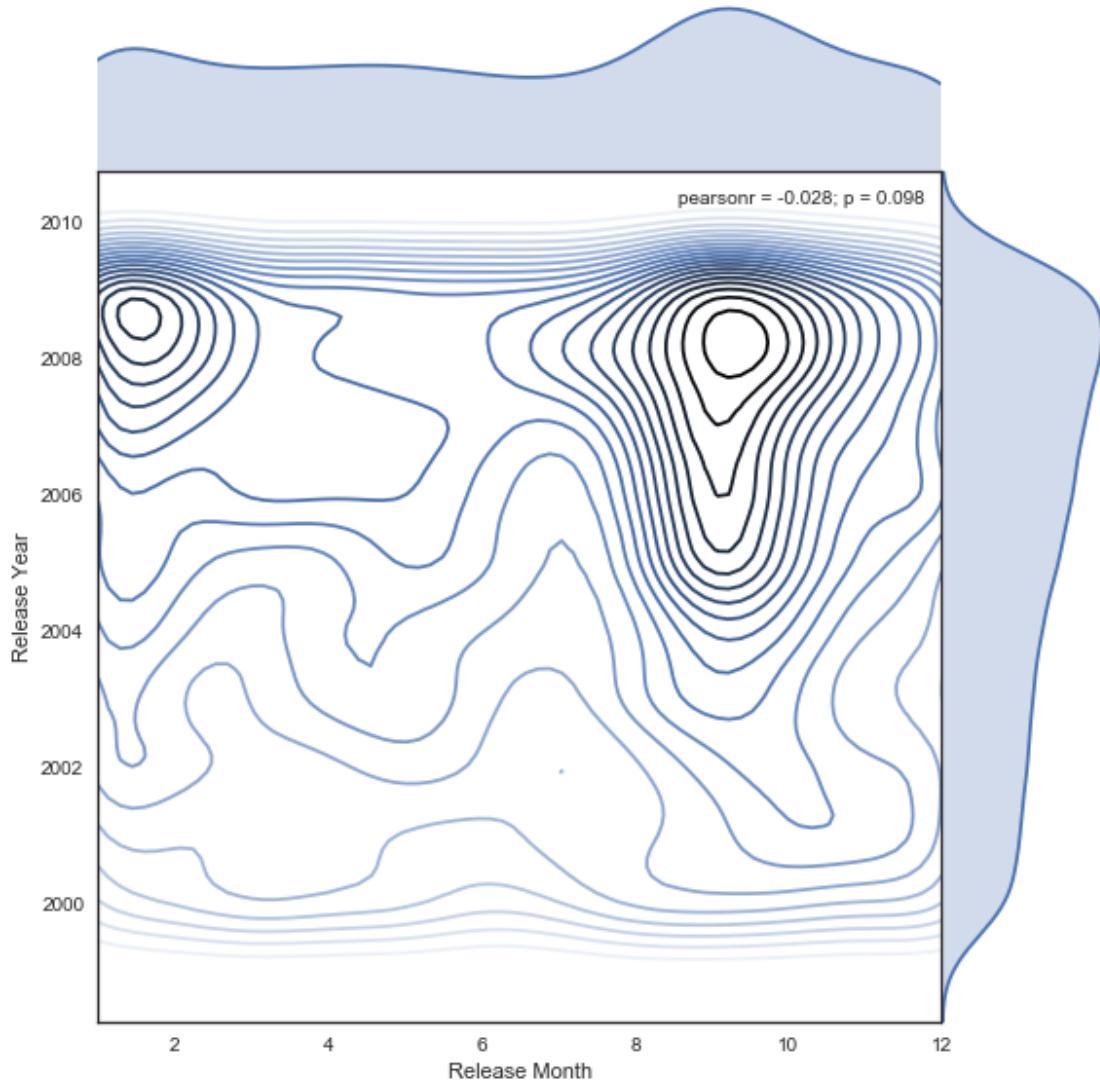
```
In [30]: # Create a contour plot for release_month vs. release_decade
for decade in ['1960s', '1970s', '1980s', '1990s', '2000s', '2010s']:
    query = 'release_decade == "'+decade+'"'
    with sns.axes_style('white'):
        ax = sns.jointplot("release_month", "release_year", df.query(query), kind='kde',
                           size=8, space=0, shade=False, n_levels=20, xlim={1,12});
        ax.ax_joint.set_xlabel('Release Month')
        ax.ax_joint.set_ylabel('Release Year')
```

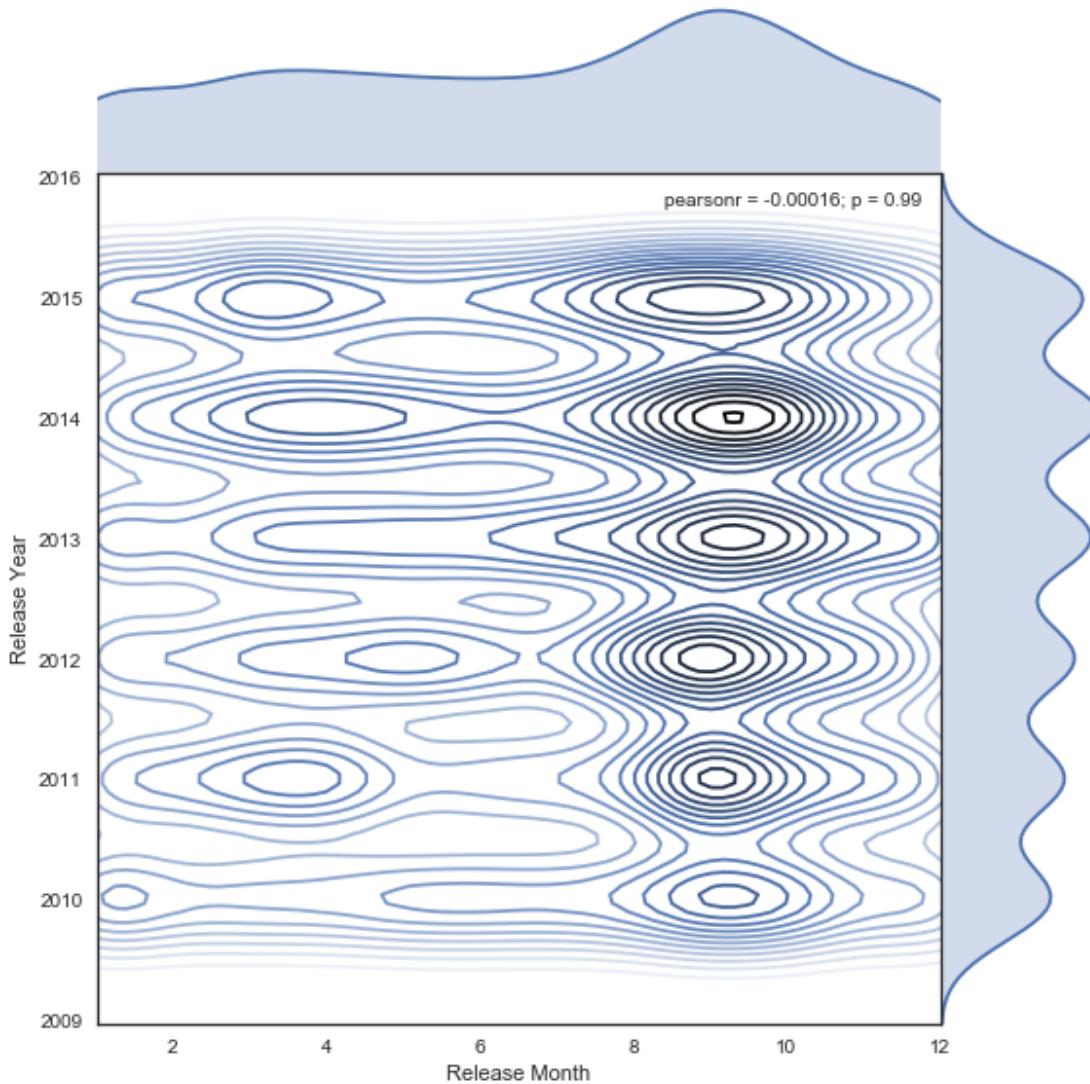












Above diagram highlights a couple of observations

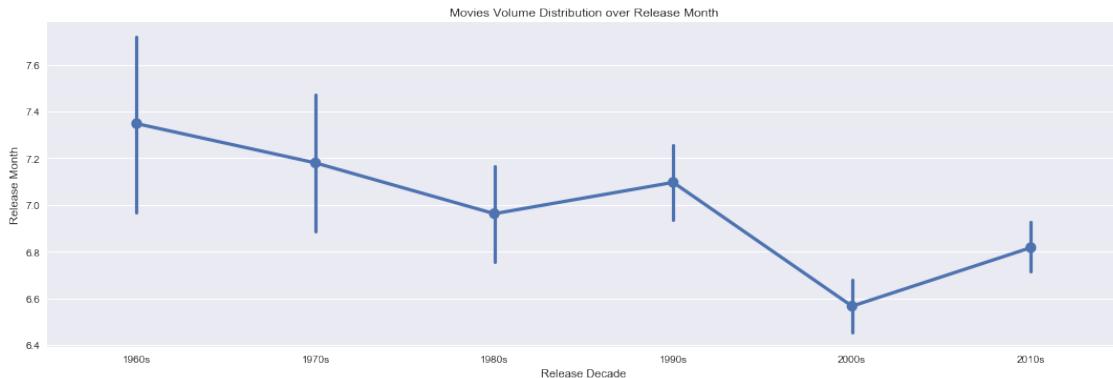
### Release Pattern

- Release pattern in recent years (2010-2015) has moved to be annual focused with the majority of movies released around summer time.
- Release pattern in earlier years (2010 and earlier) span across the whole year.
- Peak release time in earlier years (2010 and earlier) has been changing between late Winter, Spring & Summer.

**Release Volume** - The volume of released movies in the recent years (2010-2015) is significantly higher in the Summer period compared to the past years (2010 and earlier). This conclusion justifies the hilly pattern we see in the last graph).

Below diagram visualizes the last point (*recent years have narrower spread compared to earlier years*) in a more explicit way.

```
In [31]: sns.factorplot(x="release_decade", y="release_month", data=df,
                      size=5, aspect=3);
plt.title('Movies Volume Distribution over Release Month')
plt.xlabel('Release Decade')
plt.ylabel('Release Month');
```



Above diagram highlights a couple of observations - The peak of movies release has changed from first half of July to last half of June over the years. - The spread of movies release has become narrower over time.

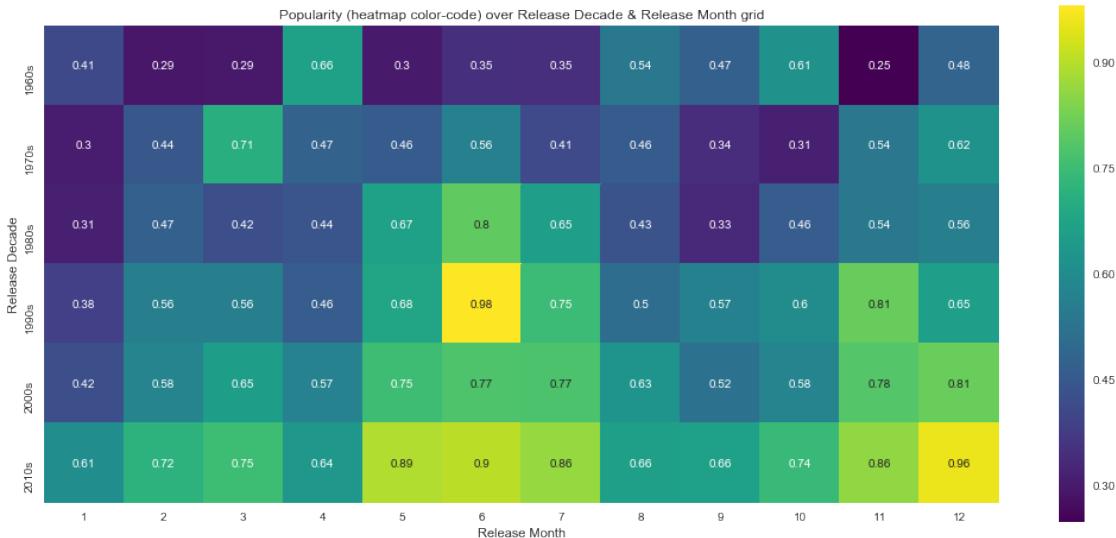
### Analyzing Popularity vs. Release time (Year/Month)

```
In [32]: df['popularity'].describe()
```

```
Out[32]: count    10834.000000
mean        0.647776
std         1.001316
min         0.000065
25%        0.208387
50%        0.384587
75%        0.715767
max        32.985763
Name: popularity, dtype: float64
```

```
In [33]: # Generate a pivot table from df breaking popularity over release_decade & release_month
result = df.pivot_table(index='release_decade', columns='release_month', values='popularity')

# Generate a heatmap using the created pivot table
plt.figure(figsize=(18,8))
sns.heatmap(result, annot=True, fmt=".2g", cmap='viridis', square=True);
plt.title('Popularity (heatmap color-code) over Release Decade & Release Month grid')
plt.xlabel('Release Month')
plt.ylabel('Release Decade');
```



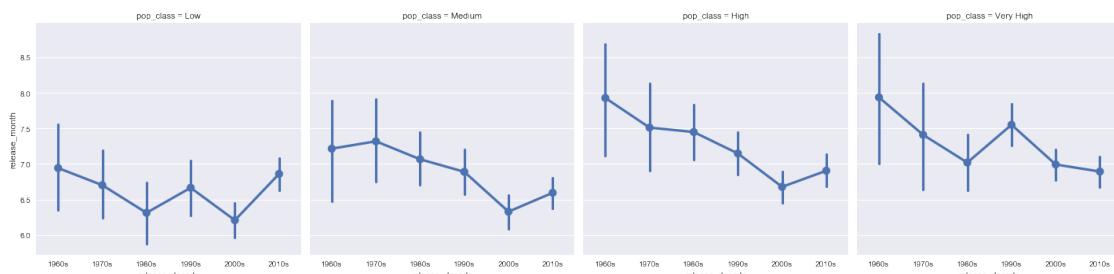
Above heat map infer a couple of observations, - Movies popularity (on average) have been increasing over the years. - Movies peak popularity tend to be around Spring (April) & Fall (October) during the 60s and 70s. That pattern has changed now to be strictly around summer (May, June, July) and Winter (November/December).

Creating a categorical column `pop_class` around popularity interquartile range 'IQR'.

```
In [34]: # Create popularity class column
bin_edges = [0.000065, 0.208387, 0.384587, 0.715767, 32.985763]
bin_names = ['Low', 'Medium', 'High', 'Very High']
df['pop_class'] = pd.cut(df['popularity'], bin_edges, labels=bin_names)
df['pop_class'].value_counts()
```

```
Out[34]: Very High    2709
          Low        2709
          High       2708
          Medium     2708
Name: pop_class, dtype: int64
```

```
In [35]: sns.factorplot(x="release_decade", y="release_month", data=df, col="pop_class",
                      size=5, aspect=1);
```



Above diagram shows that the most popular movies tend to be around a broader spread of months (June, July, August) in the 60s, That pattern has changed where the majority of most popular movies are releasing around the second half of June after the year of 2000

### Analyzing Movies Run Time

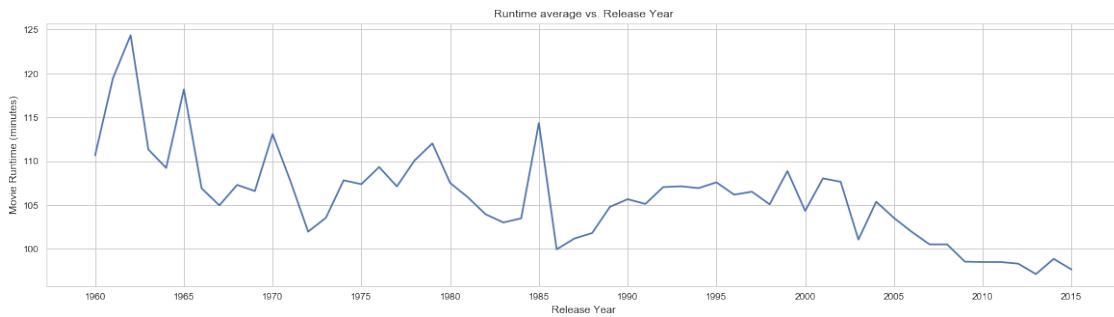
How the runtime of movies differ year to year?

In [36]: *#how the runtime of the movies differ year to year.*

```
plt.figure(figsize=(20,5))

#make the group of the data according to their release_year and find the mean related
with sns.axes_style('whitegrid'):
    df.groupby('release_year').mean()['runtime'].plot(xticks = np.arange(1960,2016,5));

#setup the figure properties
plt.title("Runtime average vs. Release Year")
plt.xlabel('Release Year')
plt.ylabel('Movie Runtime (minutes)');
```



Above diagram highlights the negative trend of **average** movie runtime over the years.

### What kind of properties associated with high-profit movies? This investigation will identify the spread in **profit** levels then focus on the potential correlation between reaching high-profit level vs. **Budget**, **Release time**, **Popularity**, **Vote Count**, **Vote Average**, and **Runtime**.

Identifying the spread in profit column

In [37]: `df['profit_adj'].describe()`

```
Out[37]: count      3.854000e+03
          mean      9.282470e+07
          std       1.940715e+08
          min      -4.139124e+08
          25%     -1.504995e+06
          50%      2.737064e+07
          75%      1.074548e+08
          max      2.750137e+09
          Name: profit_adj, dtype: float64
```

Create a profit categorical column `profit_class` breaking the data set into three classes - **Profitless** - where profit is less than 0 - **Profitable** - where profit is > 0 and less than Q3 (75% point) - **Highly Profitable** - where profit is higher than Q3 (75% point)

```
In [38]: # Create profit_class categorical column
bin_edges = [-4.139124e+08, 0, 1.074548e+08, 2.750137e+09]
bin_names = ['Profitless', 'Profitable', 'Highly Profitable']
df['profit_class'] = pd.cut(df['profit_adj'], bin_edges, labels=bin_names)
df['profit_class'].value_counts()
```

```
Out[38]: Profitable      1814
Profitless       1075
Highly Profitable    964
Name: profit_class, dtype: int64
```

```
In [39]: df['profit_class'].describe()
```

```
Out[39]: count      3853
unique        3
top      Profitable
freq      1814
Name: profit_class, dtype: object
```

Building the correlation matrix across the whole data-set

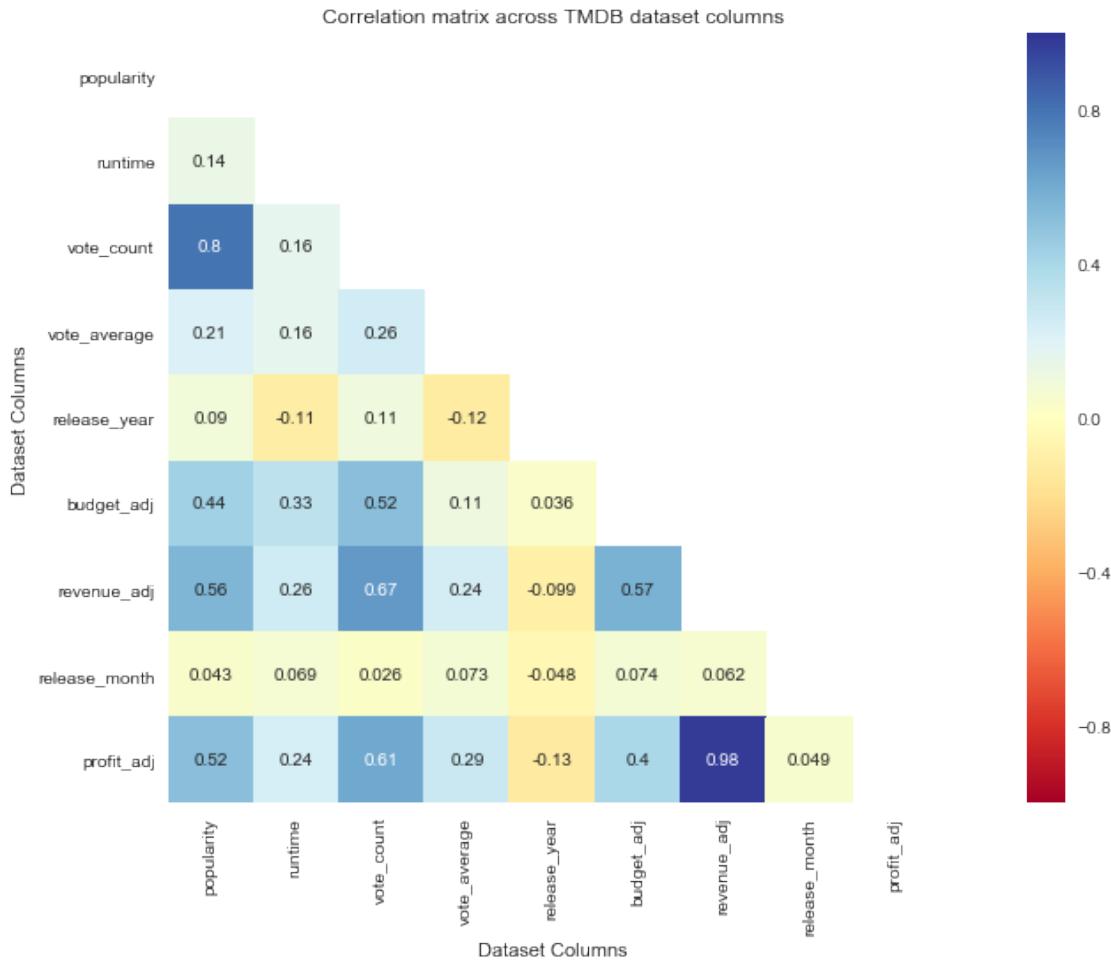
```
In [40]: #identify correlation across df & plot them in a heat map
df_corr = df.corr()

# create seaborn heatmap for df_corr
plt.figure(figsize=(18,8))
sns.set_style("whitegrid")

# generate a mask for the upper triangle
mask = np.zeros_like(df_corr, dtype=np.bool)
mask[np.triu_indices_from(mask)] = True

# plot heatmap
sns.heatmap(df_corr, annot=True, fmt=".2g", cmap='RdYlBu', mask=mask, square=True,
            center=0, vmin=-1.0, vmax=1.0);

plt.title('Correlation matrix across TMDB dataset columns')
plt.xlabel('Dataset Columns')
plt.ylabel('Dataset Columns');
```



Above matrix infers the following correlations between adjusted profit profit\_adj and - revenue\_adj with positive correlation at +0.98 - vote\_count with positive correlation at +0.61 - popularity with positive correlation at +0.52 - budget\_adj with positive correlation at +0.40 - vote\_average with positive correlation at +0.29 - runtime with positive correlation at +0.24 - release\_year with **negative** correlation at -0.13

The above matrix also infers a significant positive correlation between vote\_count & popularity

Identifying correlation across dataframe for different profit classes profit\_class

```
In [46]: # setup the plot properties
sns.set_style("whitegrid")
fig, ax = plt.subplots(1,3,sharex=True, sharey=True)
cbar_ax = fig.add_axes([.91, .12, .02, .7])

fig.suptitle('Correlation Matrix against Profit Class', fontsize = 16)

df_corr_class = {}
i = 0
```

```

for profit_level in df['profit_class'].unique():
    if profit_level != profit_level: #check if profit_level is NaN
        pass
    else:
        df_corr_class[profit_level] = df.query('profit_class == "' + profit_level + '"').corr()

    # generate a mask for the upper triangle
    mask = np.zeros_like(df_corr_class[profit_level], dtype=np.bool)
    mask[np.triu_indices_from(mask)] = True

    # plot heatmap
    sns.heatmap(df_corr_class[profit_level], annot=True, fmt=".1g", cmap='RdYlBu',
                center=0, vmin=-1.0, vmax=1.0,
                ax = ax[i], cbar=i == 0, cbar_ax=None if i else cbar_ax);
    ax[i].set_title(profit_level)

i+=1

```



Above plot do NOT infer any significant discrepancy between **Highly Profitable & Profitable** classes from correlation with other parameters prospective.

Plotting the correlation between profit\_adj and other dataset columns.

```
In [42]: # setup the plot properties
sns.set_style("whitegrid")
ax_rows, ax_columns = 2,3
```

```

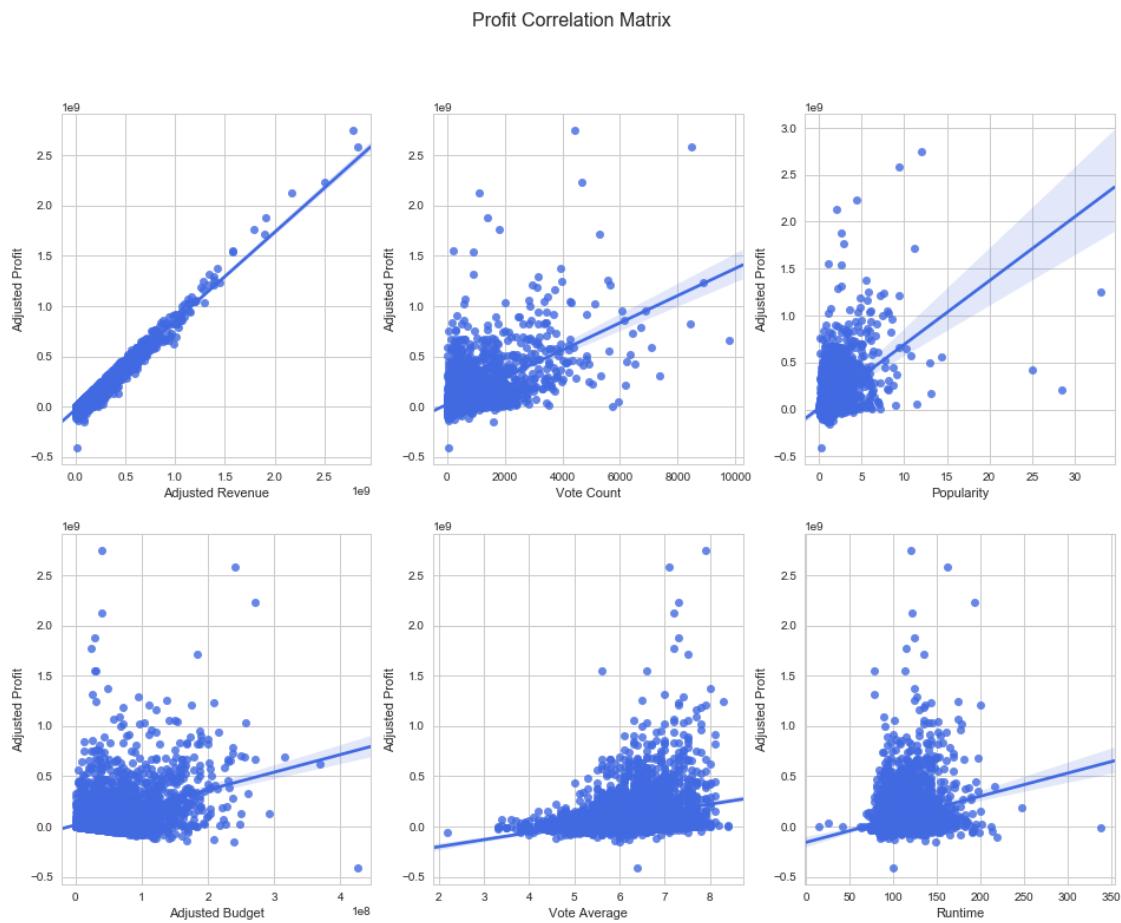
fig, ax = plt.subplots(ax_rows,ax_columns,figsize = (16,12))
fig.suptitle('Profit Correlation Matrix',fontsize = 16)

variables = [df['revenue_adj'],df['vote_count'],df['popularity'],df['budget_adj'],df['v']
variables_titles =[ 'Adjusted Revenue','Vote Count','Popularity','Adjusted Budget','Vote

for i in range(len(variables)):
    #setup subplot location
    ax_row = i//ax_columns
    ax_column = i%ax_columns

    #plot subplot with associated parameters
    sns.regplot(y=df['profit_adj'], x=variables[i], color='royalblue', ax=ax[ax_row][ax_
    ax[ax_row][ax_column].set_xlabel(variables_titles[i])
    ax[ax_row][ax_column].set_ylabel("Adjusted Profit");

```



Plotting the correlation between adjusted profit & release year

In [47]: # create a regression plot fpr revenue\_adj & budget\_adj  
`plt.figure()`

```

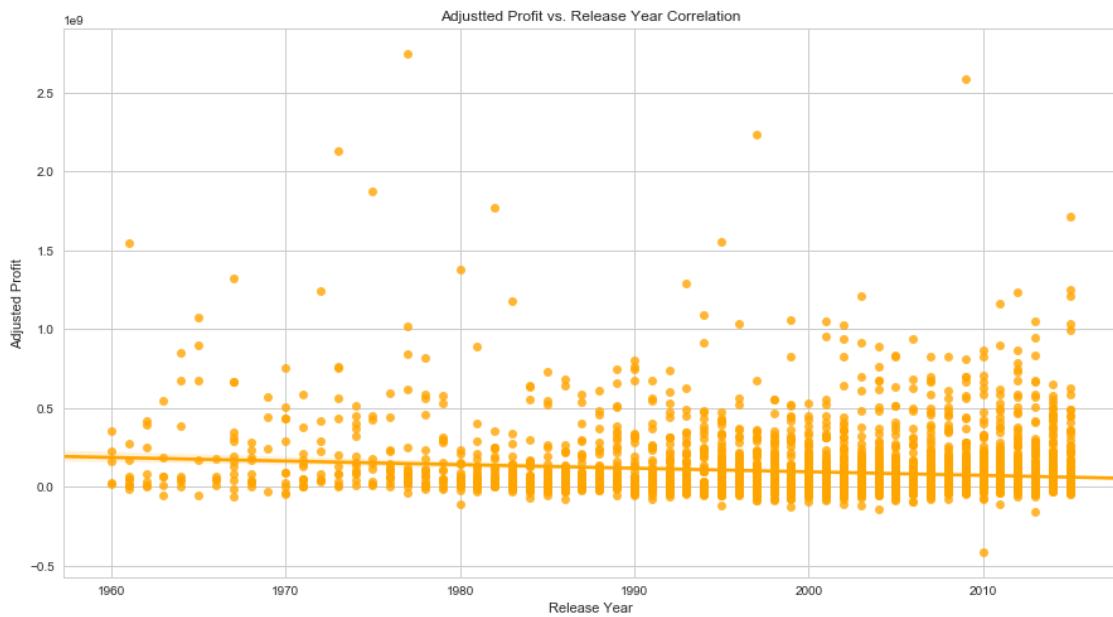
ax = sns.regplot(x=df['release_year'], y=df['profit_adj'], color='orange')

# set graph properties
sns.set(rc={'figure.figsize':(15,8)})
sns.set_style("whitegrid")
ax.set_title("Adjusted Profit vs. Release Year Correlation")
ax.set_xlabel("Release Year")
ax.set_ylabel("Adjusted Profit")

print("Adjust Profie vs. Release Year Correlation: {:.4f}".format(df_corr.loc['profit_'])

```

Adjust Profie vs. Release Year Correlation: -0.1322



Plotting correlation between Profit & other parameters **per profit\_class**.

```

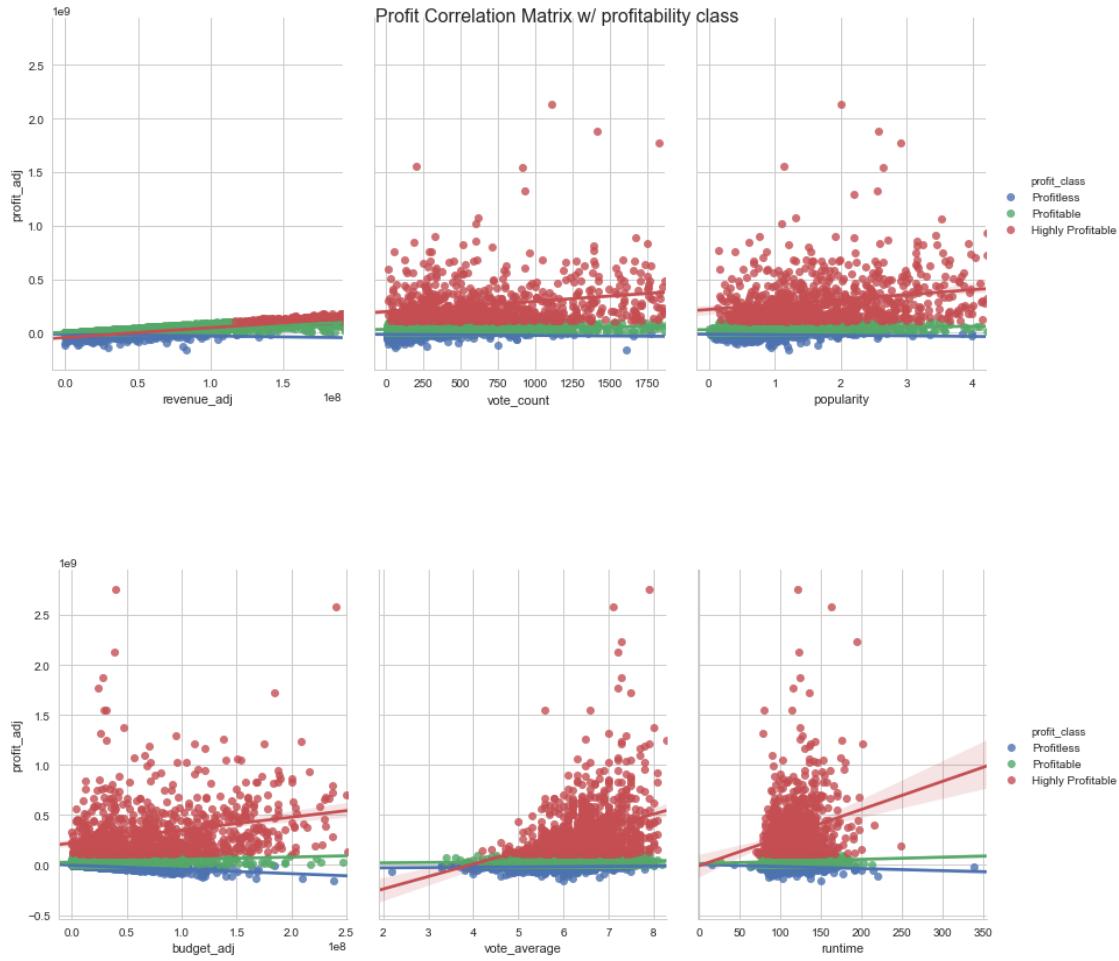
In [44]: # setup the plot properties
sns.set_style("whitegrid")

variables = ['revenue_adj','vote_count','popularity','budget_adj','vote_average','runti
variables_titles =[ 'Adjusted Revenue', 'Vote Count', 'Popularity', 'Adjusted Budget', 'Vote

#plot subplot with assoicated parameters
g = sns.pairplot(data=df, y_vars=['profit_adj'], x_vars=variables[:3], kind='reg', hue=
                  size=5, aspect=.8,);
sns.pairplot(data=df, y_vars=['profit_adj'], x_vars=variables[3:], kind='reg', hue='pro
                  size=5, aspect=.8,);

g.fig.suptitle('Profit Correlation Matrix w/ profitability class', fontsize = 16);

```



## ## Conclusions

### 0.1.3 Investigations

**Trends & patterns explored:** - The trend of release volume over the years. - The release pattern within the year. - The trend of runtime over the years.

#### Corelations evaluated:

- Popularity vs. Release Time.
- Release-month vs. Revenue & Popularity.
- Profit vs. other properties in the dataset.

### 0.1.4 Findings

- Movies production has been increased in volume significantly in the last 15 years (2000 till 2015)
- Movies tend to be released in bursts every couple of years except during the 2010s where analysis shows a focused release during summer every year.

- Release pattern in recent years (2010-2015) has moved to be annual focused with the majority of movies released around summer time.
- Release pattern in previous years (2010 and earlier) span across the whole year.
- Peak release time in earlier years (2010 and earlier) has been changing between late Winter, Spring & Summer.
- Movies peak popularity tend to be around Spring (April) & Fall (October) during the 60s and 70s. That pattern has changed now to be strictly around summer (May, June, July) and Winter (November/December).
- Runtime trend to go lower over the years.
- Analysis infers the following correlations for **Adjusted Profit**
  - Positive correlation with **Adjusted Revenue**, **Vote Count**, **Popularity**, **Adjusted Budget**, **Vote Average**, and **Runtime**.
  - Negative correlation with **Release Year**
  - Above analysis has been consistent for both Profitable & Highly Profitable classes.

### 0.1.5 Limitations

- Conclusions in this report have assumed that the provided dataset is a comprehensive set of all movies released between 1960 and 2015.

## 0.2 Submition Routine

```
In [1]: from subprocess import call
        call(['python', '-m', 'nbconvert', 'Investigate_a_Dataset.ipynb'])

Out[1]: 255
```